

# Econometrics

## Lecture 2: Simple Regression

# Recap: Foundations from Lecture 1

## Key results we will build on:

- The **conditional expectation function** (CEF)  $\mathbb{E}[Y | X]$  is the best predictor of  $Y$  given  $X$
- The **law of iterated expectations** (LIE):  $\mathbb{E}[Y] = \mathbb{E}[\mathbb{E}[Y | X]]$
- **Estimators are random variables** — properties depend on their sampling distributions
- The **central limit theorem** (CLT) justifies Normal-based inference for large samples

Today's question: how do we estimate  $\mathbb{E}[Y | X]$  from data, using a linear model?

# How Does $Y$ Relate to $X$ ?

## Motivating questions:

- How does an additional year of **education** affect **wages**?
- How does **class size** affect student **test scores**?
- How does **fertilizer use** affect crop **yield**?
- How does **job training** affect worker **productivity**?

In each case, we want to:

- ① **Quantify** the relationship between  $Y$  and  $X$
- ② **Test** whether the relationship is statistically significant
- ③ **Interpret** the result — ideally as a causal effect

# The Simple Linear Regression Model

## Simple Linear Regression

$$Y_i = \beta_0 + \beta_1 X_i + u_i, \quad i = 1, \dots, n$$

### Components:

- $Y_i$ : **dependent variable** (outcome, response, regressand)
- $X_i$ : **independent variable** (regressor, covariate, explanatory variable)
- $\beta_0$ : **intercept** parameter
- $\beta_1$ : **slope** parameter
- $u_i$ : **error term** (disturbance) — everything affecting  $Y$  besides  $X$

“Simple” = one explanatory variable. Multiple regression (Lecture 3) allows many.

# Interpreting $\beta_0$ and $\beta_1$

In the model  $Y_i = \beta_0 + \beta_1 X_i + u_i$ :

A one-unit increase in  $X$  changes  $Y$  by  $\beta_1$  units, holding  $u$  fixed. Under  $\mathbb{E}[u | X] = 0$  (SLR.4, formalized shortly):  $\partial \mathbb{E}[Y | X] / \partial X = \beta_1$ .

**Intercept**  $\beta_0$ : predicted value of  $Y$  when  $X = 0$

- Often not economically meaningful (is  $X = 0$  realistic?)

**Slope**  $\beta_1$ : the marginal effect of  $X$  on  $Y$

- **Ceteris paribus** interpretation requires assumptions on  $u$
- “Holding other factors constant” — but which factors?

# The Error Term $u_i$

The error term  $u_i$  represents **all factors** affecting  $Y_i$  besides  $X_i$ :

**In the wage equation**  $\ln(\text{wage}_i) = \beta_0 + \beta_1 \text{educ}_i + u_i$ , the error includes:

- Innate ability and intelligence
- Family background and connections
- Work experience and tenure
- Motivation and work ethic
- Measurement error in wages or education

The error term is **not** just “randomness.” It contains real, potentially important factors. Whether ordinary least squares (OLS) gives us a causal estimate depends on how  $u_i$  relates to  $X_i$ .

# Population vs. Sample Regression

## Population Regression Function (PRF)

$$\mathbb{E}[Y | X] = \beta_0 + \beta_1 X$$

- Describes the *true* relationship
- $\beta_0, \beta_1$  are unknown parameters
- We never observe this directly

## Sample Regression Function (SRF)

$$\hat{Y}_i = \hat{\beta}_0 + \hat{\beta}_1 X_i$$

- Our *estimate* from data
- $\hat{\beta}_0, \hat{\beta}_1$  computed from sample
- Changes with every new sample

The goal: use the SRF to learn about the PRF. How close is  $\hat{\beta}_1$  to  $\beta_1$ ?

# Returns to Education Revisited

From Lecture 1, the **Mincer wage equation** (Mincer, 1974):

$$\ln(\text{wage}_i) = \beta_0 + \beta_1 \text{educ}_i + u_i$$

## Interpretation:

- Because the dependent variable is  $\ln(\text{wage})$ ,  $\beta_1$  measures the **approximate percentage change** in wages per additional year of education
- If  $\hat{\beta}_1 = 0.08$ : one more year of education is associated with  $\approx 8\%$  higher wages

## The challenge:

- People with more education may also have higher ability
- If ability is in  $u_i$  and correlated with  $\text{educ}_i$ , then  $\hat{\beta}_1$  captures *both* the effect of education *and* ability
- This is the **omitted variable bias** problem (Lecture 7)

# When Can We Interpret $\beta_1$ Causally?

$\hat{\beta}_1$  always measures a **statistical association**. For it to be **causal**, we need:

The error term  $u_i$  must be **unrelated** to  $X_i$  in a specific sense:

$$\mathbb{E}[u \mid X] = 0$$

This is the **zero conditional mean** assumption — the single most important condition in this course.

## When does it fail?

- **Omitted variables:** a factor in  $u$  is correlated with  $X$
- **Reverse causality:**  $Y$  also affects  $X$
- **Selection bias:** the sample is not representative conditional on  $X$

We formalize this as Assumption SLR.4 later in this lecture.

# Roadmap for This Lecture

- ① The Simple Regression Model
- ② **Deriving the ordinary least squares (OLS) estimator**
- ③ **Assumptions for unbiasedness**
- ④ **Variance, efficiency, and the Gauss–Markov theorem**
- ⑤ **Goodness of fit:  $R^2$**
- ⑥ **Preview of inference**



**Reading:** Wooldridge (2019, Chapters 1–2)

# **Part II**

# **Deriving the OLS Estimator**

Wooldridge, Chapter 2.2

# The Idea: Best-Fitting Line

**Goal:** find the line  $\hat{Y}_i = \hat{\beta}_0 + \hat{\beta}_1 X_i$  that “best fits” the data.

## What does “best fit” mean?

- We want the line that makes the **residuals**  $\hat{u}_i = Y_i - \hat{Y}_i$  as small as possible
- We cannot just minimize  $\sum \hat{u}_i$  – positive and negative residuals cancel
- Instead, minimize the **sum of squared residuals**

**Ordinary least squares (OLS):** choose  $\hat{\beta}_0$  and  $\hat{\beta}_1$  to minimize

$$\sum_{i=1}^n \hat{u}_i^2 = \sum_{i=1}^n (Y_i - \hat{\beta}_0 - \hat{\beta}_1 X_i)^2$$

Why squared? Penalizes large deviations; yields clean, closed-form solutions.

# The OLS Objective Function

Define the sum of squared residuals as a function of  $\hat{\beta}_0$  and  $\hat{\beta}_1$ :

$$S(\hat{\beta}_0, \hat{\beta}_1) = \sum_{i=1}^n (Y_i - \hat{\beta}_0 - \hat{\beta}_1 X_i)^2$$

## OLS problem:

$$\min_{\hat{\beta}_0, \hat{\beta}_1} S(\hat{\beta}_0, \hat{\beta}_1)$$

## Strategy:

- ① Take partial derivatives with respect to  $\hat{\beta}_0$  and  $\hat{\beta}_1$
- ② Set them equal to zero (first-order conditions)
- ③ Solve the resulting system of two equations in two unknowns

$S$  is a convex quadratic in  $(\hat{\beta}_0, \hat{\beta}_1)$ , so the minimum is unique (if it exists).

# First-Order Conditions

Taking partial derivatives and setting them to zero:

**First-order condition (FOC) w.r.t.  $\hat{\beta}_0$ :**

$$\frac{\partial S}{\partial \hat{\beta}_0} = -2 \sum_{i=1}^n (Y_i - \hat{\beta}_0 - \hat{\beta}_1 X_i) = 0$$

**FOC w.r.t.  $\hat{\beta}_1$ :**

$$\frac{\partial S}{\partial \hat{\beta}_1} = -2 \sum_{i=1}^n X_i (Y_i - \hat{\beta}_0 - \hat{\beta}_1 X_i) = 0$$

Dividing by  $-2$  and rearranging gives the **normal equations**.

# The Normal Equations

Rearranging the first-order conditions:

$$\sum_{i=1}^n (Y_i - \hat{\beta}_0 - \hat{\beta}_1 X_i) = 0$$

$$\sum_{i=1}^n X_i (Y_i - \hat{\beta}_0 - \hat{\beta}_1 X_i) = 0$$

## Interpretation:

- **First equation:** residuals sum to zero:  $\sum_{i=1}^n \hat{u}_i = 0$
- **Second equation:** residuals are uncorrelated with  $X$ :  $\sum_{i=1}^n X_i \hat{u}_i = 0$

These are the **sample analogs** of  $\mathbb{E}[u] = 0$  and  $\mathbb{E}[Xu] = 0$  — the foundation of the **method of moments**.

# The OLS Estimators

Solving the normal equations yields closed-form solutions:

**OLS estimator of the slope:**

$$\hat{\beta}_1 = \frac{\sum_{i=1}^n (X_i - \bar{X})(Y_i - \bar{Y})}{\sum_{i=1}^n (X_i - \bar{X})^2}$$

**OLS estimator of the intercept:**

$$\hat{\beta}_0 = \bar{Y} - \hat{\beta}_1 \bar{X}$$

where  $\bar{X} = \frac{1}{n} \sum_{i=1}^n X_i$  and  $\bar{Y} = \frac{1}{n} \sum_{i=1}^n Y_i$  are the sample means. The OLS line always passes through the point  $(\bar{X}, \bar{Y})$ .

# Intuition for $\hat{\beta}_1$

$$\hat{\beta}_1 = \frac{\sum_{i=1}^n (X_i - \bar{X})(Y_i - \bar{Y})}{\sum_{i=1}^n (X_i - \bar{X})^2} = \frac{\text{sample covariance of } X \text{ and } Y}{\text{sample variance of } X}$$

## Reading the formula:

- **Numerator:** how  $X$  and  $Y$  move together (co-movement)
- **Denominator:** how much  $X$  varies (spread of  $X$ )
- $\hat{\beta}_1$  **scales** the co-movement by the variation in  $X$

## Connection to correlation:

$$\hat{\beta}_1 = \widehat{\text{Corr}}(X, Y) \cdot \frac{\hat{\sigma}_Y}{\hat{\sigma}_X}$$

If  $X$  and  $Y$  are positively correlated,  $\hat{\beta}_1 > 0$ . The stronger the correlation and the larger  $\hat{\sigma}_Y/\hat{\sigma}_X$ , the larger  $|\hat{\beta}_1|$ .

# Fitted Values and Residuals

## Fitted Values and Residuals

**Fitted values:**  $\hat{Y}_i = \hat{\beta}_0 + \hat{\beta}_1 X_i$

**Residuals:**  $\hat{u}_i = Y_i - \hat{Y}_i$

(predicted  $Y$  for observation  $i$ )

(prediction error for observation  $i$ )

### Key relationship:

$$Y_i = \hat{Y}_i + \hat{u}_i = \hat{\beta}_0 + \hat{\beta}_1 X_i + \hat{u}_i$$

Every observation decomposes into a **fitted part** (explained by the model) and a **residual** (unexplained).

### Important distinction:

- **Errors**  $u_i$ : unobserved, theoretical
- **Residuals**  $\hat{u}_i$ : observed, computed from the sample

# Algebraic Properties of OLS

These properties hold **by construction** — they follow from the normal equations, not from any assumptions about the data:

- ①  $\sum_{i=1}^n \hat{u}_i = 0$  (residuals sum to zero)
- ②  $\sum_{i=1}^n X_i \hat{u}_i = 0$  (residuals uncorrelated with  $X$ )
- ③  $\bar{Y} = \hat{\beta}_0 + \hat{\beta}_1 \bar{X}$  (line passes through  $(\bar{X}, \bar{Y})$ )
- ④  $\bar{\hat{Y}} = \bar{Y}$  (mean of fitted values equals mean of  $Y$ )

Properties 1–2 are the first-order conditions. Properties 3–4 follow directly.

# OLS Derivation: Summary

**OLS minimizes**  $\sum_{i=1}^n (Y_i - \hat{\beta}_0 - \hat{\beta}_1 X_i)^2$  and yields:

$$\hat{\beta}_1 = \frac{\sum_{i=1}^n (X_i - \bar{X})(Y_i - \bar{Y})}{\sum_{i=1}^n (X_i - \bar{X})^2}$$

$$\hat{\beta}_0 = \bar{Y} - \hat{\beta}_1 \bar{X}$$

**What we have so far:** a mechanical procedure to fit a line to data.

**What we still need:**

- Under what conditions does  $\hat{\beta}_1$  tell us something meaningful about  $\beta_1$ ?
- Is  $\hat{\beta}_1$  unbiased? How precise is it? How well does the line fit?

This is where assumptions come in.

# **Part III**

# **Assumptions and Properties**

Wooldridge, Chapter 2.5

# Why Do We Need Assumptions?

OLS is a **purely mechanical** procedure:

- You can always compute  $\hat{\beta}_0$  and  $\hat{\beta}_1$  from data
- OLS does not care where the data came from or what the model means
- Without assumptions,  $\hat{\beta}_1$  is just a number — not an estimate of anything

**Assumptions serve two purposes:**

- ① **Interpretation:** connect  $\hat{\beta}_1$  to the population parameter  $\beta_1$
- ② **Properties:** guarantee desirable behavior (unbiasedness, efficiency, valid inference)

The assumptions we state are *sufficient conditions*. When they fail, OLS may still work — but we can no longer guarantee it.

# Assumption SLR.1: Linear in Parameters

**SLR.1 (Linear in Parameters):** The population model is

$$Y = \beta_0 + \beta_1 X + u$$

where  $\beta_0$  and  $\beta_1$  are unknown parameters and  $u$  is the error term.

**What “linear in parameters” means:**

- $Y$  is a *linear function* of  $\beta_0$  and  $\beta_1$
- $X$  can be transformed:  $\ln(X)$ ,  $X^2$ ,  $1/X$  are all fine
- $Y = \beta_0 + \beta_1 X^2 + u$  **is** linear in parameters
- $Y = \beta_0 + \beta_1^2 X + u$  **is not** linear in parameters

“Linear regression” refers to linearity in  $\beta$ , not in  $X$ .

# Assumption SLR.2: Random Sampling

**SLR.2 (Random Sampling):** We have a random sample  $\{(X_i, Y_i) : i = 1, \dots, n\}$  from the population model.

**This means:**

- Each  $(X_i, Y_i)$  is drawn independently from the same joint distribution
- The relationship  $Y_i = \beta_0 + \beta_1 X_i + u_i$  holds for every observation
- The parameters  $\beta_0, \beta_1$  are the *same* for all  $i$

**When might this fail?**

- **Time series:** observations are typically dependent over time
- **Clustered data:** students in the same school are correlated
- **Self-selection:** only certain people respond to a survey

# Assumption SLR.3: Sample Variation in $X$

**SLR.3 (Sample Variation):** The sample values  $X_1, \dots, X_n$  are not all equal:

$$\sum_{i=1}^n (X_i - \bar{X})^2 > 0$$

## Why do we need this?

- Recall:  $\hat{\beta}_1 = \frac{\sum(X_i - \bar{X})(Y_i - \bar{Y})}{\sum(X_i - \bar{X})^2}$
- If all  $X_i$  are equal, the denominator is **zero** —  $\hat{\beta}_1$  is undefined
- **Intuition:** you cannot learn the effect of  $X$  on  $Y$  if  $X$  never varies

This is a *sample* requirement. More variation in  $X$  leads to more precise estimates — we formalize this when deriving  $\text{Var}(\hat{\beta}_1)$ .

# Assumption SLR.4: Zero Conditional Mean

**SLR.4 (Zero Conditional Mean):** The error has zero mean conditional on  $X$ :

$$\mathbb{E}[u \mid X] = 0$$

## What this says:

- For any value of  $X$ , the average of the unobserved factors is zero
- Knowing  $X$  gives **no information** about  $u$
- The population regression function is:  $\mathbb{E}[Y \mid X] = \beta_0 + \beta_1 X$

## Implications:

- $\mathbb{E}[u] = 0$  (by LIE:  $\mathbb{E}[u] = \mathbb{E}[\mathbb{E}[u \mid X]] = 0$ )
- $\text{Cov}(X, u) = 0$  (implied, but the converse is *false*)

This is the **single most important assumption** in this course. Nearly every endogeneity problem reduces to a violation of  $\mathbb{E}[u \mid X] = 0$ .

# Proving Unbiasedness of OLS

**Claim:** Under SLR.1–SLR.4,  $\mathbb{E}[\hat{\beta}_1 | \mathbf{X}] = \beta_1$ .

**Step 1:** Substitute  $Y_i = \beta_0 + \beta_1 X_i + u_i$  into  $\hat{\beta}_1$ :

$$\hat{\beta}_1 = \beta_1 + \frac{\sum_{i=1}^n (X_i - \bar{X}) u_i}{\sum_{i=1}^n (X_i - \bar{X})^2}$$

**Step 2:** Conditional expectation given  $\mathbf{X} = (X_1, \dots, X_n)$ :

$$\mathbb{E}[\hat{\beta}_1 | \mathbf{X}] = \beta_1 + \frac{\sum (X_i - \bar{X}) \mathbb{E}[u_i | \mathbf{X}]}{\sum (X_i - \bar{X})^2}$$

**Step 3:** By SLR.2 + SLR.4:  $\mathbb{E}[u_i | \mathbf{X}] = \mathbb{E}[u_i | X_i] = 0$ :

$\mathbb{E}[\hat{\beta}_1 | \mathbf{X}] = \beta_1$ . By LIE:  $\mathbb{E}[\hat{\beta}_1] = \mathbb{E}[\mathbb{E}[\hat{\beta}_1 | \mathbf{X}]] = \beta_1$ . OLS is **unbiased**.

# What $\mathbb{E}[u \mid X] = 0$ Rules Out

When  $\mathbb{E}[u \mid X] \neq 0$ , OLS is **biased**:  $\mathbb{E}[\hat{\beta}_1] \neq \beta_1$ .

## 1. Omitted variable bias — a factor in $u$ is correlated with $X$

- Education example: ability is in  $u$  and correlates with educ
- $\hat{\beta}_1$  captures education effect + ability effect  $\Rightarrow$  upward bias

## 2. Reverse causality — $Y$ also affects $X$

- Police and crime: more crime  $\rightarrow$  more police  $\rightarrow$  positive  $\text{Cov}(X, u)$

## 3. Measurement error — $X$ is measured with noise

- If we observe  $X = X^* + \varepsilon$  instead of the true  $X^*$ , the noise correlates with observed  $X$

Solutions: multiple regression (L3), IV (L8), panel data (L10).

# Assumption SLR.5: Homoskedasticity

**SLR.5 (Homoskedasticity):** The error has constant variance conditional on  $X$ :

$$\text{Var}(u \mid X) = \sigma^2$$

**Intuition:** the “spread” of  $u$  around zero does not depend on  $X$ .

**When might this fail?**

- Wage variance often *increases* with education level
- Consumption variance often *increases* with income

When  $\text{Var}(u \mid X)$  varies with  $X$ , we have **heteroskedasticity** (Lecture 6).

SLR.5 is *not* needed for unbiasedness — only for the standard variance formula and the Gauss–Markov result.

# Variance of $\hat{\beta}_1$

Under SLR.1–SLR.5, the conditional variance of  $\hat{\beta}_1$  is:

$$\text{Var}(\hat{\beta}_1 \mid \mathbf{X}) = \frac{\sigma^2}{\sum_{i=1}^n (X_i - \bar{X})^2} = \frac{\sigma^2}{\text{SST}_x}$$

where  $\text{SST}_x = \sum_{i=1}^n (X_i - \bar{X})^2$  is the total variation in  $X$ .

## What makes $\hat{\beta}_1$ more precise?

- ① **Smaller  $\sigma^2$ :** less noise in  $u \Rightarrow$  sharper estimates
- ② **Larger  $\text{SST}_x$ :** more variation in  $X \Rightarrow$  more information about the slope
- ③ **Larger  $n$ :** more data  $\Rightarrow \text{SST}_x$  grows  $\Rightarrow$  variance shrinks

Research design implication: choose  $X$  values with high spread when possible.

# Estimating $\sigma^2$ : The Standard Error of the Regression

The error variance  $\sigma^2 = \text{Var}(u \mid X)$  is unknown. We estimate it using residuals:

$$\hat{\sigma}^2 = \frac{1}{n-2} \sum_{i=1}^n \hat{u}_i^2$$

**Why  $n - 2$ ?**

- We lose 2 **degrees of freedom** from estimating  $\beta_0$  and  $\beta_1$
- The correction ensures unbiasedness:  $\mathbb{E}[\hat{\sigma}^2] = \sigma^2$

The **standard error of the regression** (SER) is:

$$\hat{\sigma} = \sqrt{\hat{\sigma}^2}$$

It measures the typical size of the OLS residuals – the “spread” of  $Y$  around the regression line, in the units of  $Y$ .

# Standard Errors of the OLS Estimators

Replacing  $\sigma^2$  with  $\hat{\sigma}^2$  gives the estimated standard errors:

**Slope:**  $se(\hat{\beta}_1) = \frac{\hat{\sigma}}{\sqrt{SST_x}}$

**Intercept:**  $se(\hat{\beta}_0) = \hat{\sigma} \cdot \sqrt{\frac{\sum_{i=1}^n X_i^2}{n \cdot SST_x}}$

**Standard errors** measure the precision of our estimates. They are used for:

- **Confidence intervals:**  $\hat{\beta}_1 \pm t_{n-2, \alpha/2} \cdot se(\hat{\beta}_1)$
- **Hypothesis tests:**  $t = \frac{\hat{\beta}_1 - \beta_{1,0}}{se(\hat{\beta}_1)}$

Full inference is covered in Lectures 3–4. For now, focus on what drives precision.

# The Gauss–Markov Theorem

**Gauss–Markov Theorem:** Under SLR.1–SLR.5, the OLS estimators  $\hat{\beta}_0$  and  $\hat{\beta}_1$  are the **best linear unbiased estimators** (BLUE).

## What “BLUE” means:

- **Best:** lowest variance among all estimators that are...
- **Linear:** ...linear functions of  $Y_1, \dots, Y_n$ , and...
- **Unbiased:** ...unbiased for  $\beta_j$

**Intuition:** among all unbiased estimators that are linear in  $Y$ , OLS has the highest precision.

**Caveat:** Gauss–Markov does *not* say OLS is the best of all estimators — only the best *linear unbiased* one. Biased or nonlinear estimators can have lower mean squared error.

# Assumption SLR.6: Normality of Errors

**SLR.6 (Normality):**  $u \mid X \sim N(0, \sigma^2)$

Combined with SLR.1–SLR.5, this gives the **classical linear model** (CLM) assumptions.

## What normality buys us:

- OLS estimators are *exactly* Normal:  $\hat{\beta}_1 \mid \mathbf{X} \sim N(\beta_1, \sigma^2/\text{SST}_x)$
- The  $t$ -statistic follows an *exact*  $t$ -distribution (not just approximately)
- Enables exact finite-sample inference (no need for large  $n$ )

## Without normality:

- OLS is still BLUE (Gauss–Markov only needs SLR.1–SLR.5)
- Inference is valid *asymptotically* via CLT (Lecture 5)
- For large  $n$ , the normality assumption matters very little

# Summary: The SLR Assumptions

- SLR.1:**  $Y = \beta_0 + \beta_1 X + u$  (linear in parameters)
- SLR.2:**  $\{(X_i, Y_i)\}_{i=1}^n$  is a random sample (random sampling)
- SLR.3:**  $\sum(X_i - \bar{X})^2 > 0$  (variation in  $X$ )
- SLR.4:**  $\mathbb{E}[u | X] = 0$  (zero conditional mean)
- SLR.5:**  $\text{Var}(u | X) = \sigma^2$  (homoskedasticity)

What you get	Assumptions needed
OLS computable	SLR.1–SLR.3
Unbiasedness	SLR.1–SLR.4
Standard variance formula	SLR.1–SLR.5
BLUE (Gauss–Markov)	SLR.1–SLR.5
Exact $t$ and $F$ tests	SLR.1–SLR.5 + SLR.6

# **Part V**

# **Measuring Fit**

Wooldridge, Chapter 2.3

# Decomposing Total Variation in $Y$

**Total variation:** how much does  $Y$  vary around its mean?

$$Y_i - \bar{Y} = \underbrace{(\hat{Y}_i - \bar{Y})}_{\text{explained}} + \underbrace{\hat{u}_i}_{\text{residual}}$$

Squaring and summing (the cross-term vanishes since  $\sum \hat{u}_i (\hat{Y}_i - \bar{Y}) = 0$ ):

$$\underbrace{\sum_{i=1}^n (Y_i - \bar{Y})^2}_{\text{TSS}} = \underbrace{\sum_{i=1}^n (\hat{Y}_i - \bar{Y})^2}_{\text{ESS}} + \underbrace{\sum_{i=1}^n \hat{u}_i^2}_{\text{RSS}}$$

- **TSS** – total sum of squares: total variation in  $Y$
- **ESS** – explained sum of squares: variation explained by the model
- **RSS** – residual sum of squares: unexplained variation

# $R^2$ : The Coefficient of Determination

$R^2$

$$R^2 = \frac{\text{ESS}}{\text{TSS}} = 1 - \frac{\text{RSS}}{\text{TSS}}$$

The **fraction of the total variation in  $Y$  explained by the regression.**

## Properties:

- $0 \leq R^2 \leq 1$  (in simple regression with an intercept)
- $R^2 = 0$ : the model explains *none* of the variation ( $\hat{\beta}_1 = 0$ )
- $R^2 = 1$ : the model explains *all* of the variation (all points on the line)
- In simple regression:  $R^2 = [\widehat{\text{Corr}}(Y, X)]^2$

# What $R^2$ Does and Does Not Tell You

## $R^2$ is useful for:

- Summarizing how well the model fits the data
- Comparing nested models (same dependent variable)

## $R^2$ does not tell you:

- Whether  $\hat{\beta}_1$  is causal
- Whether the model is correctly specified
- Whether the right variables are included
- That a “high”  $R^2$  means the model is “good”

A low  $R^2$  does not mean the regression is useless. If  $\hat{\beta}_1$  is unbiased and precisely estimated, the effect of  $X$  on  $Y$  can be well-identified even when many other factors also affect  $Y$ .

# Example: $R^2$ in the Wage Equation

Suppose we estimate:

$$\widehat{\ln(\text{wage})} = 0.584 + 0.083 \cdot \text{educ}, \quad R^2 = 0.186$$

**Interpretation of  $R^2 = 0.186$ :**

- Education explains about 19% of the variation in log wages
- The other 81% is due to experience, ability, occupation, region, etc.

**Is this “bad”? Not necessarily.**

- In cross-sectional microdata,  $R^2$  of 0.1–0.3 is common
- Humans differ in *many* ways — no single variable explains most of the variation
- $\hat{\beta}_1 = 0.083$  ( $\approx 8.3\%$  return per year) is economically significant

Illustrative; based on typical Current Population Survey wage regressions. See Wooldridge (2019, Ch. 2).

## Preview: Inference for $\beta_1$

With the SLR assumptions, we can **test hypotheses** and build **confidence intervals**:

**t-test** for  $H_0: \beta_1 = 0$  (“ $X$  has no effect on  $Y$ ”):

$$t = \frac{\hat{\beta}_1}{\text{se}(\hat{\beta}_1)} \sim t_{n-2} \quad \text{under } H_0 \text{ and SLR.1–SLR.6}$$

**95% confidence interval** for  $\beta_1$ :

$$\hat{\beta}_1 \pm t_{n-2, 0.025} \cdot \text{se}(\hat{\beta}_1)$$

Full development in Lectures 3–4. Standard errors are the bridge from estimation to inference.

# Next Time: Multiple Regression

In Lecture 3, we extend to multiple explanatory variables:

$$Y_i = \beta_0 + \beta_1 X_{1i} + \beta_2 X_{2i} + \cdots + \beta_k X_{ki} + u_i$$

## Why we need multiple regression:

- Control for confounding variables (reduce omitted variable bias)
- Isolate the effect of one variable, holding others constant
- Improve model fit and prediction

**What we will cover:** OLS in matrix form, interpreting coefficients, the Frisch–Waugh–Lovell theorem,  $R^2$  and adjusted  $R^2$ .

**Reading:** Wooldridge (2019, Chapter 3)

# Key Takeaways

- ➊ The simple regression model  $Y = \beta_0 + \beta_1 X + u$  relates an outcome to one explanatory variable
- ➋ **OLS** minimizes  $\sum \hat{u}_i^2$  and yields  $\hat{\beta}_1 = \frac{\text{sample cov. of } X, Y}{\text{sample var. of } X}$
- ➌ Under SLR.1–SLR.4, OLS is **unbiased**:  $\mathbb{E}[\hat{\beta}_1] = \beta_1$
- ➍ The **zero conditional mean** assumption  $\mathbb{E}[u \mid X] = 0$  is crucial for causal interpretation
- ➎ Under SLR.1–SLR.5, OLS is **BLUE** (Gauss–Markov theorem)
- ➏  $R^2$  measures fit but says nothing about causality

## References I

- Mincer, J. (1974). *Schooling, Experience, and Earnings*. Columbia University Press, New York.
- Wooldridge, J. M. (2019). *Introductory Econometrics: A Modern Approach*. Cengage Learning, Boston, MA, 7th edition.